

Assignment 3

- Dumbvm is a very limited virtual memory system
 - A full TLB leads to a kernel panic
 - Text segment is not read-only
 - Uses fixed segmentation
 - External fragmentation
 - Never reuses physical memory
 - Requires restarting the OS after each test
- Assignment 3 fixes these problems!

TLB Replacement

- VM related exceptions are handled by *vm_fault*
- *vm_fault* performs address translation and loads the virtual address to physical address mapping into the TLB.
 - Iterates through the TLB to find an unused/invalid entry
 - Overwrites the unused entry with the virtual to physical address mapping required by the instruction that generated the TLB exception
- If the TLB is full, call *tlb_random* to write the entry into a random TLB slot
 - That's it for TLB replacement!
 - Make sure that virtual page fields in the TLB are unique

Read-Only Text Segment

- Currently, TLB entries are loaded with **TLBLO_DIRTY** on
 - Pages are therefore read and writeable
- Text segment should be read-only
 - Load TLB entries for the text segment with TLBLO_DIRTY off
 - `elo &= ~TLBLO_DIRTY;`
- Determine the segment of the fault address by looking at the vbase and vtop addresses

Read-Only Text Segment

- Unfortunately, this will cause `load_elf` to throw a `VM_FAULT_READONLY` exception
 - It is trying to write to a memory location that is read-only
- We must instead load TLB entries with `TLBLO_DIRTY` on until `load_elf` has completed.
 - Consider adding a flag to **`struct addrspace`** to indicate whether or not `load_elf` has completed
 - When `load_elf` completes, flush the TLB, and ensure that all future TLB entries for the text segment has `TLBLO_DIRTY` off

Read-Only Text Segment

- Writing to read-only memory address will lead to a `VM_FAULT_READONLY` exception
 - This will currently cause a kernel panic
- Instead of panicing, your VM system should kill the process
 - Modify `kill_curthread` to kill the current process
 - Very similar to `sys__exit`. However, the exit code/status will be different
 - Consider which signal number this will trigger (look at the beginning of `kill_curthread`)

Managing Memory

Physical memory

0x0

A large, empty rectangle with a thin blue border, representing a block of physical memory. It is positioned centrally below the title and above the address label.

Managing Memory

Physical memory

During bootstrap, the kernel allocates memory by calling `getppages`, which in turn calls `ram_stealmem(pages)`.

`ram_stealmem` just allocates pages without providing any mechanism to release pages (see `free_kpages`)

We want to manage our own memory after bootstrap


0x0



Memory for bootstrap

Managing Memory

Physical memory



Memory for bootstrap


In `vm_bootstrap`, call `ram_getsize` to get the remaining physical memory in the system. Do not call `ram_stealmem` again!

Logically partition the memory into fixed size frames. Each frame is `PAGE_SIZE` bytes.

Keep track of the status of each frame (core-map data structure).

Managing Memory

Physical memory



Core-map
Memory for bootstrap

Where should we store the core-map data structure?

Store it in the start of the memory returned by `ram_getsize`. The frames should start after the core-map data structure.

The core-map should track which frames are used and available. It should also keep track of contiguous memory allocations (why?)

Alloc and Free

- `alloc_kpages(int npages):`
 - Allocate frames for both `kmalloc` and address spaces
 - Frames need to be contiguous
- `free_kpages(vaddr_t addr):`
 - Free pages allocated with `alloc_kpages`
 - We don't specify how many pages we need to free. It should free the same number of pages that was allocated.
 - Update core-map to make frames available after `free_kpages`
- Consider adding some information in the core-map to help determine the number of pages that need to be freed
 - E.g. If 4 contiguous frames were allocated using `alloc_kpages`, then store 4 in the core-map entry for the start of the four frames

Page Tables

- In order to avoid external fragmentation, we need to introduce paging
- New VM system combines segmentation with paging
- Three segments:
 - Text (read-only)
 - Data
 - Stack
- Create a page table for each segment
 - Each page table entry should include the frame number

Page Tables

- In dumbvm, struct addrspace has the following fields:
 - vaddr_t as_vbase1
 - paddr_t as_pbase1
 - size_t as_npages1
 - vaddr_t as_vbase2
 - paddr_t as_pbase2
 - size_t as_npages2
 - paddr_t as_stackpbase
- With segmentation and paging, what fields should we have instead?
 - Replace pbase with page table

Page Tables

- `as_create`:
 - Initialize address space data structures
- `as_define_region`:
 - A region is essentially a segment
 - Allocate (`kmalloc`) and initialize the page table for the specified segment
 - Size of the segment is a parameter of `as_define_region`
 - Because we perform preloading, segment size will never grow!
 - Size of the page table is based on the segment size
 - Setup the read/write permissions for this segment
 - Optionally have permissions per page

Page Tables

- `as_prepare_load`:
 - Pre-allocate frames for each page in the segments
 - Frames do not need to be contiguous
 - Allocate each frame one at a time
 - Need to create a page table for the stack
 - `as_define_region` not called for the stack segment

Page Tables

- `as_copy`:
 - Call *as_create* to create the address space
 - Create segments based on old address space
 - Allocate frames for the segments
 - `memcpy` frames from the old address space to the frames in the new address space
- `as_destroy`:
 - Call `free_kpages` on the frames for each segment
 - `kfree` the page tables

User Address/Kernel Virtual Address/ Physical Address

- Remember that you are always working with virtual addresses
 - Only use physical addresses when loading entries in the TLB
 - Virtual addresses are converted either by the TLB or by the MMU directly
- Addresses below 0x80000000 are user space addresses that are TLB mapped
- Addresses between 0x80000000 and 0xa0000000 are kernel virtual addresses that are converted by the MMU directly
 - Kernel virtual address – 0x80000000 = physical address
- kmalloc always returns a kernel virtual address. Do not use kmalloc to allocate frames