

# Intrinsic Plasticity and Batch Normalisation

Nolan Peter Shaw

2019-03-15


## Agenda

- Introduce Intrinsic Plasticity (IP)
- Discuss the biological and computational benefits of IP
- Introduce batch normalisation (BN)
- Outline BN implementation
- Demonstrate the relation between IP and BN


## Section I: Intrinsic Plasticity

Computational Neuroscience


## Synaptic vs Intrinsic Plasticity in the Brain




## Synaptic vs Intrinsic Plasticity in the Brain



## Synaptic vs Intrinsic Plasticity in the Brain




## Intrinsic Plasticity




## Biological Benefits

- Human brain consumes calories
- Cost of a 0/1 in an ANN is identical



Triesch, Jochen. (2005). A Gradient Rule for the Plasticity of a Neuron's Intrinsic Excitability. Artificial Neural Networks: Biological Inspirations ICANN 2005.

## Computational Benefits



Bell AJ, Sejnowski TJ (November 1995). An information-maximization approach to blind separation and blind deconvolution. Neural Computation

## Computational Benefits

### Synaptic Plasticity

- Learn weights w.r.t. an error signal

- Minimise loss on some task

### Intrinsic Plasticity

- Learn gains and biases w.r.t. local statistics

- Maximise information potential

## Implementation in ANNs

### Biology

- Sensitivity

- Threshold

### Artificial

- Gain (Horizontal stretch)

- Bias (Horizontal translation)

### Original Activation Function

### Becomes

$$y = \theta(x) \longrightarrow y = \theta(\alpha * x + k)$$

## Implementation in ANNs

Original Activation Function

Becomes

$$y = \theta(x) \longrightarrow y = \theta(\alpha * x + k)$$

Super simple (YAY!)

## Implementation in ANNs

Update rules

Gain:

Bias:

$$\Delta\alpha = \frac{1}{\alpha} - 2 * \mathbf{E}[\mathbf{xy}]$$

$$\Delta k = -2 * \mathbf{E}[\mathbf{y}]$$

- Note that these update rules are still being studied and that there may be update rules that are better suited to learning

## Issues

- Unstable
- $\mathbf{E}[\mathbf{uy}]$  may be ill-suited for adjusting the sensitivity/gain
- May homogenise inputs too much
- Competes with error-based learning of synaptic weights

## Section II: Batch Normalisation

Machine Learning

## The Problem

- The “shape” of inputs to a layer may be radically different from one input to the next
- This slows down learning as hidden layers are required to learn representations and distributions as well as perform computation

## The Problem

- The “shape” of inputs to a layer may be radically different from one input to the next
- This slows down learning as hidden layers are required to learn representations and distributions as well as perform computation

## The Solution

- Normalise all inputs w.r.t. their distribution
  - (infeasible to do for an entire dataset so treat each batch as a sample of the population)
- De-normalise w.r.t. error
  - (prevents homogeneity and preserves computational properties of neurons)

## Visualising Batch Normalisation





Image courtesy of: <https://zaffnet.github.io/batch-normalization>

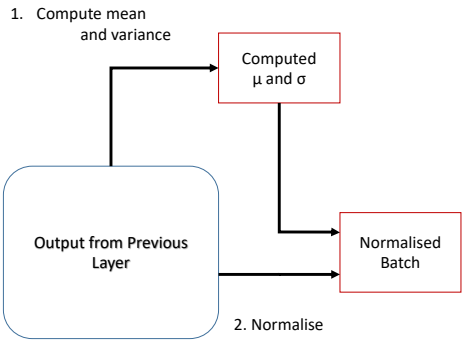
## Implementing Batch Normalisation

<b>Input:</b> Values of $x$ over a mini-batch: $\mathcal{B} = \{x_1 \dots x_m\}$ ;	
Parameters to be learned: $\gamma, \beta$	
<b>Output:</b> $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$	
$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i$	// mini-batch mean
$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2$	// mini-batch variance
$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}}$	// normalize
$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i)$	// scale and shift

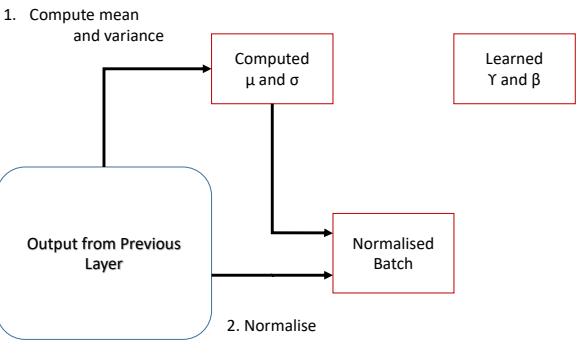
# Step-by-step Walkthrough



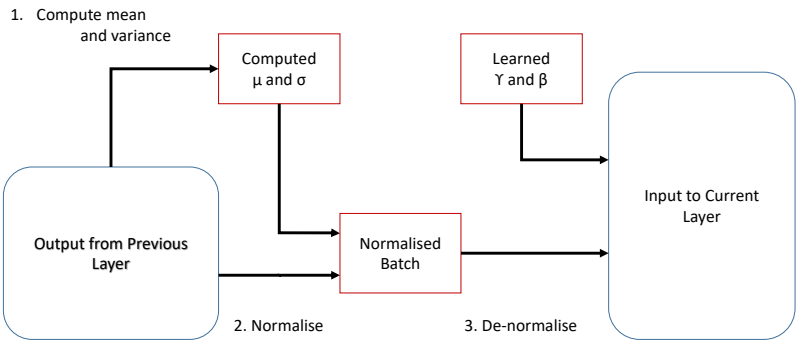
# Step-by-step Walkthrough



# Step-by-step Walkthrough



# Step-by-step Walkthrough



## Section III: Unifying IP and BN

(or: How I Stopped Caring About Big Data and Learned to Love the Brain)

Equivalence of the two models

Intrinsic Plasticity:

$$y = \theta(\alpha * x + k)$$

Equivalence of the two models

Intrinsic Plasticity:

$$y = \theta(\gamma * (\alpha * x + k) + \beta)$$

Equivalence of the two models

Intrinsic Plasticity:

$$y = \theta(\gamma * (\alpha * x + k) + \beta)$$

Batch Normalisation:

$$y = \theta(\gamma * \left( \frac{x - \mu}{\sqrt{\sigma^2 + \epsilon}} \right) + \beta)$$

## Equivalence of the two models

Intrinsic Plasticity:

$$y = \theta(\gamma * (\alpha * x + k) + \beta)$$

Batch Normalisation:

$$y = \theta\left(\gamma * \left(\frac{1}{\sqrt{\sigma^2 + \epsilon}} * x + \frac{-\mu}{\sqrt{\sigma^2 + \epsilon}}\right) + \beta\right)$$

## Equivalence of the two models

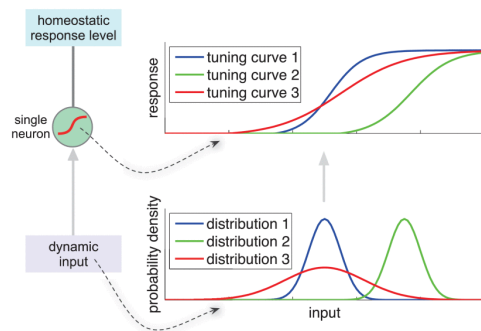
Intrinsic Plasticity:

$$y = \theta(\gamma * (\alpha * x + k) + \beta)$$

Batch Normalisation:

$$y = \theta\left(\gamma * \left(\frac{1}{\sqrt{\sigma^2 + \epsilon}} * x + \frac{-\mu}{\sqrt{\sigma^2 + \epsilon}}\right) + \beta\right)$$

## Relation to the Vanishing Gradient Problem



Thank you!